

NEX3

STRATEGIC ADVISORY FOR THE AI & WEB3 ECONOMY

A NEX3 FIELD GUIDE

The Ultimate Guide to **Free AI**

Replace \$300+/Month AI Subscriptions with Open Source
Never Pay for AI Again

Chat

Code

Image

Video

Voice

PREPARED BY NEX3 INC. · NEX3.XYZ

Why Open-Source AI, Why Now

If you're paying for more than one AI subscription, you've probably felt it: the bill keeps growing. A ChatGPT Plus subscription here, a Claude Pro plan there, Midjourney for images, GitHub Copilot for code, ElevenLabs for voice — it adds up fast, and most of these tools overlap in what they actually do for you.

Over the last two years, open-source AI has closed the gap with paid tools faster than almost anyone expected. Models like DeepSeek R1, Qwen 3, Llama 3.3, and Gemma 3 now compete directly with commercial chatbots on reasoning and writing quality. Image generators like FLUX.1 rival Midjourney. Local coding assistants like Aider and Cline can pair-program with you the same way Copilot does — except everything runs on hardware you already own.

What "Open Source AI" Actually Means Here

- **Open-weight models** — AI models whose trained weights you can download and run yourself (Llama, Qwen, Gemma, DeepSeek, FLUX, and others). Licensing varies by model, so check each model's license before commercial use.
- **Open-source software** — the applications that run those models: Ollama, Open WebUI, ComfyUI, ChatterUI, and similar tools, whose source code is publicly available.

BEGINNER TIP: You don't need to understand how neural networks work to use any of this. If you can install an app and copy-paste a terminal command, you can run every tool in this guide.

The Case for Running AI Locally

Benefit	What it means in practice
Cost	Pay once for hardware you may already own, then run unlimited queries with no per-token or monthly fee.
Privacy	Your prompts, documents, and code never leave your machine, which matters for sensitive business or personal data.
No rate limits	No "message caps," no waiting for your quota to reset.
Offline access	Local models work on a plane, in a basement server room, or anywhere without internet.
Full control	You choose the exact model version, fine-tune it, or swap it out any time.

WARNING: Local AI isn't free of trade-offs. You'll need a reasonably capable computer (see Chapter 3), and the very largest frontier models still tend to be the paid, cloud-hosted ones. This guide focuses on tools that get you 80–95% of the way there for \$0.

Monthly Cost Comparison

Paid Tool	Category	Typical Monthly Cost
ChatGPT Plus	Chat / reasoning	\$20
Claude Pro	Chat / reasoning	\$20
Gemini Advanced	Chat / reasoning	\$20
Midjourney	Image generation	\$30
Perplexity Pro	Search + chat	\$20
GitHub Copilot	Coding assistant	\$10–39
Runway	Video generation	\$35–95
ElevenLabs	Voice / TTS	\$22–99
Total (typical stack)		\$300–500+ / month

\$300–500 / month → \$0 / month

Paid AI stack → Open-source stack

Paid AI vs. Free Alternatives

Paid Tool	Purpose	Price	Best Free Alternatives	Diff.	Qual.
ChatGPT Plus	General chat / reasoning	\$20/mo	Llama 3.3, DeepSeek R1, Qwen 3, Gemma 3	Easy	★★★★☆
Claude Pro	Reasoning, long context, writing	\$20/mo	DeepSeek R1, Qwen 3	Easy	★★★★☆
Gemini Advanced	General chat / multimodal	\$20/mo	Gemma 3	Easy	★★★★☆
GitHub Copilot	Coding assistant	\$10–39/mo	Continue.dev, Cline, Aider, OpenHands	Medium	★★★★☆
Midjourney	Image generation	\$30/mo	FLUX.1-dev, SDXL, ComfyUI	Medium	★★★★★
Perplexity	AI-powered search	\$20/mo	Open WebUI + SearXNG	Medium	★★★★☆
Runway	Video generation	\$35–95/mo	Wan 2.2, CogVideoX, Stable Video Diffusion	Hard	★★★★☆
ElevenLabs	Text-to-speech, voice cloning	\$22–99/mo	Kokoro TTS, Piper, XTTS	Easy	★★★★☆
Notion AI	Notes + AI workspace	\$10/mo	AppFlowy, AFFiNE, Obsidian	Easy	★★★★☆
Jasper	AI writing / marketing copy	\$39–125/mo	LibreChat, AnythingLLM	Medium	★★★★☆

PRO TIP: You don't need to replace every tool at once. Start with Ollama + Open WebUI (Chapters 4–5) since that single combo replaces ChatGPT Plus, Claude Pro, and Gemini Advanced for most day-to-day use.

Meet the Free Alternatives

Chat & Reasoning Models — **Llama 3.3** (Meta): the most broadly supported open-weight model family. **DeepSeek R1**: specializes in step-by-step reasoning for math, logic, and coding. **Qwen 3** (Alibaba): strong multilingual performance and long context windows. **Gemma 3** (Google DeepMind): lightweight and efficient on modest hardware.

Coding Assistants — **Continue.dev**: drop-in Copilot replacement for VS Code / JetBrains. **Cline**: autonomous VS Code agent for multi-file edits. **Aider**: terminal-first AI pair programmer integrated with git. **OpenHands**: fully autonomous coding agent for multi-step engineering tasks.

Image Generation — **FLUX.1** (Black Forest Labs): leads open-source image models on prompt accuracy and photorealism. **SDXL** (Stability AI): most mature model with a huge library of community fine-tunes. **ComfyUI**: the node-based interface most serious users run FLUX and SDXL through.

Search, Video & Voice — **Open WebUI + SearXNG**: approximates Perplexity's search-plus-chat experience. **Wan 2.2** (Alibaba): unusually coherent motion for open-source video. **CogVideoX**: lighter-weight video model. **Kokoro TTS**: lightweight, natural-sounding TTS on CPU. **Piper**: extremely fast, ideal for real-time use. **XTTS**: voice cloning from a short reference clip.

Notes & Workspace — **AppFlowy**: open-source Notion alternative. **AFFiNE**: privacy-first Notion + Miro hybrid. **Obsidian**: local-first markdown notes with a huge plugin ecosystem. **LibreChat**: self-hosted ChatGPT-style interface. **AnythingLLM**: all-in-one local chat and document-Q&A app.

CHAPTER 03

Before You Install: Hardware, OS & Setup Basics

Minimum and Recommended Hardware

Use Case	Minimum	Recommended
Text chat (7–8B models)	16GB RAM, any modern CPU	16GB RAM + 8GB VRAM GPU
Text chat (30B+ models)	32GB RAM	24GB+ VRAM GPU (RTX 3090/4090)
Image generation (SDXL/FLUX)	8GB VRAM	12–24GB VRAM
Video generation	16GB VRAM	24GB+ VRAM
Voice (STT/TTS)	Any modern CPU	Any modern CPU (GPU optional)

BEGINNER TIP: No dedicated GPU? You can still run smaller models (Llama 3.3 8B, Gemma 3 4B/12B, Qwen 3 8B) comfortably on CPU with 16GB of RAM. Apple Silicon Macs (M1–M4) are especially efficient thanks to unified memory.

General Installation Flow



macOS: Install Homebrew first. Apple Silicon (M-series) is recommended — unified memory lets large models run without a discrete GPU. Docker Desktop is required for containerized tools.

Windows: Install WSL2 first (`wsl --install` in admin PowerShell). NVIDIA GPU users should install the latest driver plus CUDA Toolkit. Docker Desktop with WSL2 backend is standard.

Linux: Most tools ship native installers or Docker images — Linux is usually the smoothest OS for this whole stack. Install NVIDIA drivers + `nvidia-container-toolkit` for Docker GPU acceleration.

WARNING: Only download installers from the official project websites or GitHub release pages listed in this guide. Never use third-party "mirror" sites for AI model weights or installers — mirrors are a common malware vector.

Common Installation Times

Tool	Estimated Time	Difficulty
Ollama	5 minutes	Easy
Open WebUI (Docker)	10 minutes	Easy
Continue.dev / Cline (VS Code extension)	5 minutes	Easy
ComfyUI	20–40 minutes	Medium
Video generation (Wan 2.2 / CogVideoX)	45–90 minutes	Hard

CHAPTER 04

Ollama: The Complete Guide

Ollama is the easiest way to run open-weight language models locally. It handles model downloading, quantization, and serving through a single command-line tool — think of it as "Docker for AI models." It downloads and manages models from its built-in library, runs an inference server at `http://localhost:11434`, and exposes a REST API that other tools connect to.

Installing Ollama

macOS: `brew install ollama` | **Windows:** Download `OllamaSetup.exe` from `ollama.com/download` | **Linux:** `curl -fsSL https://ollama.com/install.sh | sh`

Core Commands

Command	What it does
<code>ollama pull llama3.3</code>	Downloads a model to your machine
<code>ollama run llama3.3</code>	Starts an interactive chat session
<code>ollama list</code>	Shows all models you've downloaded
<code>ollama ps</code>	Shows which models are currently loaded in memory
<code>ollama rm llama3.3</code>	Deletes a model to free disk space
<code>ollama show llama3.3</code>	Displays model details (size, parameters, template)

PRO TIP: Use quantized versions (tags like `:q4_K_M`) for a good balance of speed and quality on limited hardware — they run noticeably faster with only a small quality trade-off.

Common Errors

Error	Fix
"connection refused" on port 11434	The Ollama service isn't running — start it from the app or run <code>ollama serve</code> .

Out of memory / model won't load	Pick a smaller model or a more aggressive quantization tag.
Slow responses on CPU	Expected for large models without a GPU — try a smaller model size.

NEX3 INC. — VANCOUVER, BC

Installing All of This Yourself Feels Like a Full-Time Job.

You just read four chapters' worth of terminals, Docker commands, VRAM math, and licensing footnotes. That's not a knock on you — it's genuinely a lot to stand up on your own, especially while running a business.

- Drowning in AI hype, unsure where to even start?
- Burning weeks evaluating tools that don't fit your workflow?
- Sales and ops data piling up with zero actionable insight?
- Web3 and blockchain still feel like a different language?

ONE FRAMEWORK. THREE MOVES.

- 1** We map your AI + Web3 opportunity into a 90-day plan you can actually ship — no fluff, no theory.
- 2** Hands-on sessions with a senior consultant. We build working artifacts alongside you, not slide decks.
- 3** You deploy it in your business, then expand from a working playbook — not a blank page.

"A founder told us their free discovery call turned a vague idea into a clear, actionable plan — and helped them secure a \$5,000 Vancouver grant soon after."

Skip the trial-and-error. Let's build it together.

www.nex3.xyz/contact

NEX3 INC. — VANCOUVER, BC · STRATEGIC ADVISORY FOR THE AI & WEB3 ECONOMY

NEX3

NEX3.XYZ